

# quantium

2023-03-25

## Solution template for Task 1

This file is a solution template for the Task 1 of the Quantum Virtual Internship. It will walk you through the analysis, providing the scaffolding for your solution with gaps left for you to fill in yourself. Page 1 20200128\_InsideSherpa\_Task1\_DraftSolutions - Template (1).Rmd Look for comments that say “over to you” for places where you need to add your own code! Often, there will be hints about what to do or what function to use in the text leading up to a code block - if you need a bit of extra help on how to use a function, the internet has many excellent resources on R coding, which you can find using your favourite search engine. ## Load required libraries and datasets Note that you will need to install these libraries if you have never used these before.

```
#### Example code to install packages
install.packages("data.table")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.2'
## (as 'lib' is unspecified)
```

```
install.packages("ggplot2")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.2'
## (as 'lib' is unspecified)
```

```
install.packages("ggmosaic")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.2'
## (as 'lib' is unspecified)
```

```
install.packages("readr")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.2'
## (as 'lib' is unspecified)
```

```
#### Load required libraries
```

```
library(data.table)
```

```
library(ggplot2)
```

```
library(ggmosaic)
```

```
library(readr)
```

```
#### Point the filePath to where you have downloaded the datasets to and
```

```
#### assign the data files to data.tables
```

```
# over to you! fill in the path to your working directory. If you are on a Windows machine, you will ne
filePath <- "/cloud/project/"
```

```
transactionData <- fread(paste0(filePath,"QVI_transaction_data.csv"))
```

```
customerData <- fread(paste0(filePath,"QVI_purchase_behaviour.csv"))
```

## Exploratory data analysis

The first step in any analysis is to first understand the data. Let's take a look at each of the datasets provided. ### Examining transaction data We can use `str()` to look at the format of each column and see a sample

of the data. As we have read in the dataset as a `data.table` object, we can also run `transactionData` in the console to see a sample of the data or use `head(transactionData)` to look at the first 10 rows. Let's check if columns we would expect to be numeric are in numeric form and date columns are in date format.

```
#### Examine transaction data
```

```
# Over to you! Examine the data using one or more of the methods described above.
head(transactionData)
```

```
##      DATE STORE_NBR LYLTY_CARD_NBR TXN_ID PROD_NBR
## 1: 43390         1         1000      1         5
## 2: 43599         1         1307     348        66
## 3: 43605         1         1343     383        61
## 4: 43329         2         2373     974        69
## 5: 43330         2         2426    1038       108
## 6: 43604         4         4074    2982        57
##
##              PROD_NAME PROD_QTY TOT_SALES
## 1:   Natural Chip      Compny SeaSalt175g      2      6.0
## 2:              CCs Nacho Cheese      175g      3      6.3
## 3:   Smiths Crinkle Cut  Chips Chicken 170g      2      2.9
## 4:   Smiths Chip Thinly  S/Cream&Onion 175g      5     15.0
## 5: Kettle Tortilla ChpsHny&Jlpno Chili 150g      3     13.8
## 6: Old El Paso Salsa   Dip Tomato Mild 300g      1      5.1
```

```
str(transactionData)
```

```
## Classes 'data.table' and 'data.frame': 264836 obs. of 8 variables:
## $ DATE : int 43390 43599 43605 43329 43330 43604 43601 43601 43332 43330 ...
## $ STORE_NBR : int 1 1 1 2 2 4 4 4 5 7 ...
## $ LYLTY_CARD_NBR: int 1000 1307 1343 2373 2426 4074 4149 4196 5026 7150 ...
## $ TXN_ID : int 1 348 383 974 1038 2982 3333 3539 4525 6900 ...
## $ PROD_NBR : int 5 66 61 69 108 57 16 24 42 52 ...
## $ PROD_NAME : chr "Natural Chip Compny SeaSalt175g" "CCs Nacho Cheese 175g"
##               "Smiths Crinkle Cut Chips Chicken 170g" "Smiths Chip Thinly S/Cream&Onion 175g"
## ...
## $ PROD_QTY : int 2 3 2 5 3 1 1 1 1 2 ...
## $ TOT_SALES : num 6 6.3 2.9 15 13.8 5.1 5.7 3.6 3.9 7.2 ...
## - attr(*, ".internal.selfref")=<externalptr>
```

We can see that the date column is in an integer format. Let's change this to a date format.

```
#### Convert DATE column to a date format
```

```
#### A quick search online tells us that CSV and Excel integer dates begin on 30 Dec 1899
```

```
transactionData$DATE <- as.Date(transactionData$DATE, origin = "1899-12-30")
str(transactionData)
```

```
## Classes 'data.table' and 'data.frame': 264836 obs. of 8 variables:
## $ DATE : Date, format: "2018-10-17" "2019-05-14" ...
## $ STORE_NBR : int 1 1 1 2 2 4 4 4 5 7 ...
## $ LYLTY_CARD_NBR: int 1000 1307 1343 2373 2426 4074 4149 4196 5026 7150 ...
## $ TXN_ID : int 1 348 383 974 1038 2982 3333 3539 4525 6900 ...
## $ PROD_NBR : int 5 66 61 69 108 57 16 24 42 52 ...
## $ PROD_NAME : chr "Natural Chip Compny SeaSalt175g" "CCs Nacho Cheese 175g"
##               "Smiths Crinkle Cut Chips Chicken 170g" "Smiths Chip Thinly S/Cream&Onion 175g"
## ...
## $ PROD_QTY : int 2 3 2 5 3 1 1 1 1 2 ...
## $ TOT_SALES : num 6 6.3 2.9 15 13.8 5.1 5.7 3.6 3.9 7.2 ...
```

```
## - attr(*, ".internal.selfref")=<externalptr>
```

We should check that we are looking at the right products by examining PROD\_NAME.

```
#### Examine PROD_NAME
# Over to you! Generate a summary of the PROD_NAME column.
summary(transactionData$PROD_NAME)
```

```
##      Length      Class      Mode
## 264836 character character
```

Looks like we are definitely looking at potato chips but how can we check that these are all chips? We can do some basic text analysis by summarising the individual words in the product name.

```
#### Examine the words in PROD_NAME to see if there are any incorrect entries
#### such as products that are not chips
productWords <- data.table(unlist(strsplit(unique(transactionData[, PROD_NAME]), "
")))
setnames(productWords, 'words')
head(productWords)
```

```
##                                words
## 1: Natural Chip      Compny SeaSalt175g
## 2:                CCs Nacho Cheese  175g
## 3: Smiths Crinkle Cut Chips Chicken 170g
## 4: Smiths Chip Thinly S/Cream&Onion 175g
## 5: Kettle Tortilla ChpsHny&Jlpno Chili 150g
## 6: Old El Paso Salsa  Dip Tomato Mild 300g
```

As we are only interested in words that will tell us if the product is chips or not, let's remove all words with digits and special characters such as '&' from our set of product words. We can do this using `grepl()`.

```
# Over to you! Remove digits, and special characters, and then sort the distinct words by frequency of
#### Removing digits
```

```
#### Removing special characters
productWords$words <- gsub("[^[:alpha:][:space:]]", "", productWords$words)
head(productWords)
```

```
##                                words
## 1: Natural Chip      Compny SeaSaltg
## 2:                CCs Nacho Cheese  g
## 3: Smiths Crinkle Cut Chips Chicken g
## 4: Smiths Chip Thinly SCreamOnion g
## 5: Kettle Tortilla ChpsHnyJlpno Chili g
## 6: Old El Paso Salsa  Dip Tomato Mild g
```

```
#### Let's look at the most common words by counting the number of times a word appears and
#### sorting them by this frequency in order of highest to lowest frequency
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:data.table':
##
##      between, first, last

## The following objects are masked from 'package:stats':
##
```

```
##      filter, lag
## The following objects are masked from 'package:base':
##
##      intersect, setdiff, setequal, union
# Group the data by words and count the frequency of occurrence
word_freq <- productWords %>%
  group_by(words) %>%
  summarize(count = n())
# Sort the data in descending order of frequency
word_freq <- word_freq %>%
  arrange(desc(count))
# View the top 10 most frequent words
head(word_freq, 10)
```

```
## # A tibble: 10 x 2
##   words                                count
##   <chr>                                <int>
## 1 Burger Rings g                        1
## 2 CCs Nacho Cheese g                    1
## 3 CCs Original g                        1
## 4 CCs Tasty Cheese g                    1
## 5 Cheetos Chs Bacon Balls g            1
## 6 Cheetos Puffs g                       1
## 7 Cheezels Cheese Box g                 1
## 8 Cheezels Cheese g                     1
## 9 Cobs Popd Sea Salt Chips g            1
## 10 Cobs Popd Sour Crm Chives Chips g    1
```

There are salsa products in the dataset but we are only interested in the chips category, so let's remove these.

```
#### Remove salsa products
transactionData[, SALSA := grepl("salsa", tolower(PROD_NAME))]
transactionData <- transactionData[SALSA == FALSE, ][, SALSA := NULL]
```


Next, we can use `summary()` to check summary statistics such as mean, min and max values for each feature to see if there are any obvious outliers in the data and if there are any nulls in any of the columns (NA's : number of nulls will appear in the output if there are any nulls).

```
#### Summarise the data to check for nulls and possible outliers
# Over to you!
#is.na(transactionData)
summary(transactionData)
```


```
##      DATE      STORE_NBR  LYLTY_CARD_NBR  TXN_ID
## Min.   :2018-07-01  Min.   : 1.0  Min.   : 1000  Min.   :    1
## 1st Qu.:2018-09-30  1st Qu.: 70.0  1st Qu.: 70015  1st Qu.: 67569
## Median :2018-12-30  Median :130.0  Median : 130367  Median : 135183
## Mean   :2018-12-30  Mean   :135.1  Mean   : 135531  Mean   : 135131
## 3rd Qu.:2019-03-31  3rd Qu.:203.0  3rd Qu.: 203084  3rd Qu.: 202654
## Max.   :2019-06-30  Max.   :272.0  Max.   :2373711  Max.   :2415841
##  PROD_NBR  PROD_NAME  PROD_QTY  TOT_SALES
## Min.   : 1.00  Length:246742  Min.   : 1.000  Min.   : 1.700
## 1st Qu.: 26.00  Class :character  1st Qu.: 2.000  1st Qu.: 5.800
## Median : 53.00  Mode  :character  Median : 2.000  Median : 7.400
## Mean   : 56.35              Mean   : 1.908  Mean   : 7.321
```

```
## 3rd Qu.: 87.00          3rd Qu.: 2.000    3rd Qu.: 8.800
## Max.    :114.00        Max.    :200.000    Max.    :650.000
```

```
# Create a boxplot to check for outliers
boxplot(transactionData$TOT_SALES)
```



```
boxplot(transactionData$PROD_QTY)
```



There are no nulls in the columns but product quantity appears to have an outlier which we should investigate further. Let's investigate further the case where 200 packets of chips are bought in one transaction.

```
#### Filter the dataset to find the outlier
# Over to you! Use a filter to examine the transactions in question.
transactionData_filtered <- transactionData[transactionData$PROD_QTY == 200, ]
```

There are two transactions where 200 packets of chips are bought in one transaction and both of these transactions were by the same customer.

```
#### Let's see if the customer has had other transactions
# Over to you! Use a filter to see what other transactions that customer made.
transactionData_filtered
```

```
##          DATE STORE_NBR LYLTY_CARD_NBR TXN_ID PROD_NBR
## 1: 2018-08-19      226      226000 226201      4
```

```
## 2: 2019-05-20      226      226000 226210      4
##                PROD_NAME PROD_QTY TOT_SALES
## 1: Dorito Corn Chp Supreme 380g      200      650
## 2: Dorito Corn Chp Supreme 380g      200      650

transactionData_filtered1 <- transactionData[transactionData$LYLTY_CARD_NBR == 226000, ]
transactionData_filtered1
```

```
##          DATE STORE_NBR LYLTY_CARD_NBR TXN_ID PROD_NBR
## 1: 2018-08-19      226      226000 226201      4
## 2: 2019-05-20      226      226000 226210      4
##                PROD_NAME PROD_QTY TOT_SALES
## 1: Dorito Corn Chp Supreme 380g      200      650
## 2: Dorito Corn Chp Supreme 380g      200      650
```

It looks like this customer has only had the two transactions over the year and is not an ordinary retail customer. The customer might be buying chips for commercial purposes instead. We'll remove this loyalty card number from further analysis.

```
#### Filter out the customer based on the loyalty card number
# Over to you!
#### Re-examine transaction data
# Over to you!
# Remove rows where LYLTY_CARD_NBR is equal to 226000
transactionData <- transactionData[transactionData$LYLTY_CARD_NBR != 226000, ]
# View the modified data frame
str(transactionData)
```

```
## Classes 'data.table' and 'data.frame': 246740 obs. of 8 variables:
## $ DATE : Date, format: "2018-10-17" "2019-05-14" ...
## $ STORE_NBR : int 1 1 1 2 2 4 4 5 7 7 ...
## $ LYLTY_CARD_NBR: int 1000 1307 1343 2373 2426 4149 4196 5026 7150 7215 ...
## $ TXN_ID : int 1 348 383 974 1038 3333 3539 4525 6900 7176 ...
## $ PROD_NBR : int 5 66 61 69 108 16 24 42 52 16 ...
## $ PROD_NAME : chr "Natural Chip Compny SeaSalt175g" "CCs Nacho Cheese 175g"
## "Smiths Crinkle Cut Chips Chicken 170g" "Smiths Chip Thinly S/Cream&Onion 175g"
## ...
## $ PROD_QTY : int 2 3 2 5 3 1 1 1 2 1 ...
## $ TOT_SALES : num 6 6.3 2.9 15 13.8 5.7 3.6 3.9 7.2 5.7 ...
## - attr(*, ".internal.selfref")=<externalptr>
```

That's better. Now, let's look at the number of transaction lines over time to see if there are any obvious data issues such as missing data.

```
#### Count the number of transactions by date
# Over to you! Create a summary of transaction count by date.
# Convert date column to a Date object
transactionData$DATE <- as.Date(transactionData$DATE, format = "%Y-%m-%d")

# Aggregate transaction count by date
transactions_by_date <- aggregate(x = list(count = transactionData$TXN_ID),
                                by = list(date = transactionData$DATE),
                                FUN = length)

# Print the summary
print(transactions_by_date)
```

##		date	count
## 1		2018-07-01	663
## 2		2018-07-02	650
## 3		2018-07-03	674
## 4		2018-07-04	669
## 5		2018-07-05	660
## 6		2018-07-06	711
## 7		2018-07-07	695
## 8		2018-07-08	653
## 9		2018-07-09	692
## 10		2018-07-10	650
## 11		2018-07-11	701
## 12		2018-07-12	717
## 13		2018-07-13	727
## 14		2018-07-14	661
## 15		2018-07-15	712
## 16		2018-07-16	678
## 17		2018-07-17	694
## 18		2018-07-18	689
## 19		2018-07-19	637
## 20		2018-07-20	684
## 21		2018-07-21	683
## 22		2018-07-22	673
## 23		2018-07-23	673
## 24		2018-07-24	648
## 25		2018-07-25	674
## 26		2018-07-26	672
## 27		2018-07-27	697
## 28		2018-07-28	640
## 29		2018-07-29	659
## 30		2018-07-30	692
## 31		2018-07-31	688
## 32		2018-08-01	680
## 33		2018-08-02	669
## 34		2018-08-03	662
## 35		2018-08-04	665
## 36		2018-08-05	705
## 37		2018-08-06	706
## 38		2018-08-07	668
## 39		2018-08-08	695
## 40		2018-08-09	652
## 41		2018-08-10	675
## 42		2018-08-11	678
## 43		2018-08-12	642
## 44		2018-08-13	703
## 45		2018-08-14	702
## 46		2018-08-15	702
## 47		2018-08-16	690
## 48		2018-08-17	663
## 49		2018-08-18	683
## 50		2018-08-19	670
## 51		2018-08-20	644
## 52		2018-08-21	653
## 53		2018-08-22	689

##	54	2018-08-23	696
##	55	2018-08-24	647
##	56	2018-08-25	657
##	57	2018-08-26	685
##	58	2018-08-27	670
##	59	2018-08-28	636
##	60	2018-08-29	666
##	61	2018-08-30	653
##	62	2018-08-31	658
##	63	2018-09-01	687
##	64	2018-09-02	671
##	65	2018-09-03	661
##	66	2018-09-04	718
##	67	2018-09-05	685
##	68	2018-09-06	745
##	69	2018-09-07	663
##	70	2018-09-08	666
##	71	2018-09-09	705
##	72	2018-09-10	645
##	73	2018-09-11	647
##	74	2018-09-12	661
##	75	2018-09-13	646
##	76	2018-09-14	688
##	77	2018-09-15	636
##	78	2018-09-16	669
##	79	2018-09-17	660
##	80	2018-09-18	717
##	81	2018-09-19	670
##	82	2018-09-20	656
##	83	2018-09-21	699
##	84	2018-09-22	609
##	85	2018-09-23	738
##	86	2018-09-24	672
##	87	2018-09-25	729
##	88	2018-09-26	652
##	89	2018-09-27	632
##	90	2018-09-28	694
##	91	2018-09-29	671
##	92	2018-09-30	704
##	93	2018-10-01	662
##	94	2018-10-02	650
##	95	2018-10-03	658
##	96	2018-10-04	684
##	97	2018-10-05	651
##	98	2018-10-06	702
##	99	2018-10-07	644
##	100	2018-10-08	676
##	101	2018-10-09	724
##	102	2018-10-10	700
##	103	2018-10-11	706
##	104	2018-10-12	658
##	105	2018-10-13	663
##	106	2018-10-14	636
##	107	2018-10-15	674



##	108	2018-10-16	675
##	109	2018-10-17	682
##	110	2018-10-18	611
##	111	2018-10-19	699
##	112	2018-10-20	679
##	113	2018-10-21	677
##	114	2018-10-22	684
##	115	2018-10-23	659
##	116	2018-10-24	672
##	117	2018-10-25	655
##	118	2018-10-26	716
##	119	2018-10-27	643
##	120	2018-10-28	649
##	121	2018-10-29	666
##	122	2018-10-30	665
##	123	2018-10-31	652
##	124	2018-11-01	695
##	125	2018-11-02	670
##	126	2018-11-03	680
##	127	2018-11-04	697
##	128	2018-11-05	642
##	129	2018-11-06	673
##	130	2018-11-07	679
##	131	2018-11-08	662
##	132	2018-11-09	710
##	133	2018-11-10	713
##	134	2018-11-11	731
##	135	2018-11-12	678
##	136	2018-11-13	653
##	137	2018-11-14	681
##	138	2018-11-15	689
##	139	2018-11-16	679
##	140	2018-11-17	701
##	141	2018-11-18	690
##	142	2018-11-19	722
##	143	2018-11-20	732
##	144	2018-11-21	651
##	145	2018-11-22	626
##	146	2018-11-23	702
##	147	2018-11-24	670
##	148	2018-11-25	610
##	149	2018-11-26	642
##	150	2018-11-27	680
##	151	2018-11-28	640
##	152	2018-11-29	685
##	153	2018-11-30	670
##	154	2018-12-01	675
##	155	2018-12-02	655
##	156	2018-12-03	677
##	157	2018-12-04	666
##	158	2018-12-05	660
##	159	2018-12-06	645
##	160	2018-12-07	672
##	161	2018-12-08	622

##	162	2018-12-09	659
##	163	2018-12-10	664
##	164	2018-12-11	686
##	165	2018-12-12	624
##	166	2018-12-13	668
##	167	2018-12-14	697
##	168	2018-12-15	671
##	169	2018-12-16	709
##	170	2018-12-17	729
##	171	2018-12-18	799
##	172	2018-12-19	839
##	173	2018-12-20	808
##	174	2018-12-21	781
##	175	2018-12-22	840
##	176	2018-12-23	853
##	177	2018-12-24	865
##	178	2018-12-26	700
##	179	2018-12-27	690
##	180	2018-12-28	669
##	181	2018-12-29	666
##	182	2018-12-30	686
##	183	2018-12-31	650
##	184	2019-01-01	634
##	185	2019-01-02	674
##	186	2019-01-03	637
##	187	2019-01-04	704
##	188	2019-01-05	636
##	189	2019-01-06	673
##	190	2019-01-07	668
##	191	2019-01-08	669
##	192	2019-01-09	686
##	193	2019-01-10	685
##	194	2019-01-11	631
##	195	2019-01-12	687
##	196	2019-01-13	628
##	197	2019-01-14	663
##	198	2019-01-15	657
##	199	2019-01-16	674
##	200	2019-01-17	677
##	201	2019-01-18	658
##	202	2019-01-19	696
##	203	2019-01-20	683
##	204	2019-01-21	637
##	205	2019-01-22	689
##	206	2019-01-23	647
##	207	2019-01-24	619
##	208	2019-01-25	671
##	209	2019-01-26	672
##	210	2019-01-27	648
##	211	2019-01-28	661
##	212	2019-01-29	667
##	213	2019-01-30	689
##	214	2019-01-31	690
##	215	2019-02-01	708

##	216	2019-02-02	692
##	217	2019-02-03	690
##	218	2019-02-04	659
##	219	2019-02-05	691
##	220	2019-02-06	666
##	221	2019-02-07	663
##	222	2019-02-08	714
##	223	2019-02-09	671
##	224	2019-02-10	697
##	225	2019-02-11	683
##	226	2019-02-12	684
##	227	2019-02-13	693
##	228	2019-02-14	664
##	229	2019-02-15	667
##	230	2019-02-16	670
##	231	2019-02-17	682
##	232	2019-02-18	619
##	233	2019-02-19	664
##	234	2019-02-20	695
##	235	2019-02-21	646
##	236	2019-02-22	692
##	237	2019-02-23	689
##	238	2019-02-24	682
##	239	2019-02-25	693
##	240	2019-02-26	662
##	241	2019-02-27	687
##	242	2019-02-28	682
##	243	2019-03-01	670
##	244	2019-03-02	677
##	245	2019-03-03	674
##	246	2019-03-04	670
##	247	2019-03-05	724
##	248	2019-03-06	661
##	249	2019-03-07	658
##	250	2019-03-08	701
##	251	2019-03-09	637
##	252	2019-03-10	651
##	253	2019-03-11	690
##	254	2019-03-12	711
##	255	2019-03-13	702
##	256	2019-03-14	640
##	257	2019-03-15	724
##	258	2019-03-16	664
##	259	2019-03-17	715
##	260	2019-03-18	644
##	261	2019-03-19	688
##	262	2019-03-20	692
##	263	2019-03-21	672
##	264	2019-03-22	725
##	265	2019-03-23	680
##	266	2019-03-24	710
##	267	2019-03-25	688
##	268	2019-03-26	689
##	269	2019-03-27	660

##	270	2019-03-28	677
##	271	2019-03-29	699
##	272	2019-03-30	684
##	273	2019-03-31	647
##	274	2019-04-01	635
##	275	2019-04-02	700
##	276	2019-04-03	718
##	277	2019-04-04	709
##	278	2019-04-05	664
##	279	2019-04-06	681
##	280	2019-04-07	640
##	281	2019-04-08	650
##	282	2019-04-09	646
##	283	2019-04-10	699
##	284	2019-04-11	632
##	285	2019-04-12	672
##	286	2019-04-13	664
##	287	2019-04-14	685
##	288	2019-04-15	651
##	289	2019-04-16	674
##	290	2019-04-17	649
##	291	2019-04-18	667
##	292	2019-04-19	655
##	293	2019-04-20	738
##	294	2019-04-21	710
##	295	2019-04-22	671
##	296	2019-04-23	663
##	297	2019-04-24	702
##	298	2019-04-25	701
##	299	2019-04-26	684
##	300	2019-04-27	667
##	301	2019-04-28	677
##	302	2019-04-29	697
##	303	2019-04-30	680
##	304	2019-05-01	643
##	305	2019-05-02	667
##	306	2019-05-03	657
##	307	2019-05-04	630
##	308	2019-05-05	680
##	309	2019-05-06	707
##	310	2019-05-07	667
##	311	2019-05-08	698
##	312	2019-05-09	668
##	313	2019-05-10	665
##	314	2019-05-11	679
##	315	2019-05-12	687
##	316	2019-05-13	638
##	317	2019-05-14	705
##	318	2019-05-15	632
##	319	2019-05-16	664
##	320	2019-05-17	652
##	321	2019-05-18	626
##	322	2019-05-19	730
##	323	2019-05-20	707

```
## 324 2019-05-21    671
## 325 2019-05-22    687
## 326 2019-05-23    633
## 327 2019-05-24    691
## 328 2019-05-25    705
## 329 2019-05-26    648
## 330 2019-05-27    665
## 331 2019-05-28    683
## 332 2019-05-29    714
## 333 2019-05-30    669
## 334 2019-05-31    664
## 335 2019-06-01    682
## 336 2019-06-02    662
## 337 2019-06-03    656
## 338 2019-06-04    637
## 339 2019-06-05    680
## 340 2019-06-06    700
## 341 2019-06-07    762
## 342 2019-06-08    699
## 343 2019-06-09    718
## 344 2019-06-10    676
## 345 2019-06-11    634
## 346 2019-06-12    709
## 347 2019-06-13    607
## 348 2019-06-14    743
## 349 2019-06-15    724
## 350 2019-06-16    690
## 351 2019-06-17    658
## 352 2019-06-18    639
## 353 2019-06-19    662
## 354 2019-06-20    698
## 355 2019-06-21    716
## 356 2019-06-22    643
## 357 2019-06-23    653
## 358 2019-06-24    612
## 359 2019-06-25    696
## 360 2019-06-26    657
## 361 2019-06-27    669
## 362 2019-06-28    673
## 363 2019-06-29    703
## 364 2019-06-30    704
```

There's only 364 rows, meaning only 364 dates which indicates a missing date. Let's create a sequence of dates from 1 Jul 2018 to 30 Jun 2019 and use this to create a chart of number of transactions over time to find the missing date.

```
#### Create a sequence of dates and join this the count of transactions by date
# Over to you - create a column of dates that includes every day from 1 Jul 2018 to 30 Jun 2019, and jo
# Generate a sequence of dates from 1 Jul 2018 to 30 Jun 2019
date_seq <- seq(from = as.Date("2018-07-01"), to = as.Date("2019-06-30"), by = "day")

# Create a data frame with the date sequence
date_df <- data.frame(DATE = date_seq)

# Join the date_df to transaction_by_date to fill in the missing days
```

```

transactions_by_day <- merge(date_df, transactions_by_date, all = TRUE)

library(dplyr)

# Rename the "date" column to "transaction_date" in transaction_by_day
transactions_by_day <- transactions_by_day %>% rename(transaction_date = date)

# Merge transaction_by_day with date_df to fill in the missing dates
transactions_by_day_filled <- merge(date_df, transactions_by_day, all.x = TRUE)


# Print the summary of transaction count by date with missing days filled in
head(transactions_by_day_filled)

##          DATE transaction_date count
## 1 2018-07-01      2018-07-01    663
## 2 2018-07-01      2018-07-02    650
## 3 2018-07-01      2018-07-04    669
## 4 2018-07-01      2018-07-05    660
## 5 2018-07-01      2018-07-07    695
## 6 2018-07-01      2018-07-10    650

transactions_by_day_filled <- transactions_by_day_filled %>% select(-transaction_date)

#### Setting plot themes to format graphs
theme_set(theme_bw())
theme_update(plot.title = element_text(hjust = 0.5))
#### Plot transactions over time
ggplot(transactions_by_day_filled, aes(x = DATE, y = count)) +
  geom_line() +
  labs(x = "Day", y = "Number of transactions", title = "Transactions over time") +
  scale_x_date(breaks = "1 month") +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5))

```

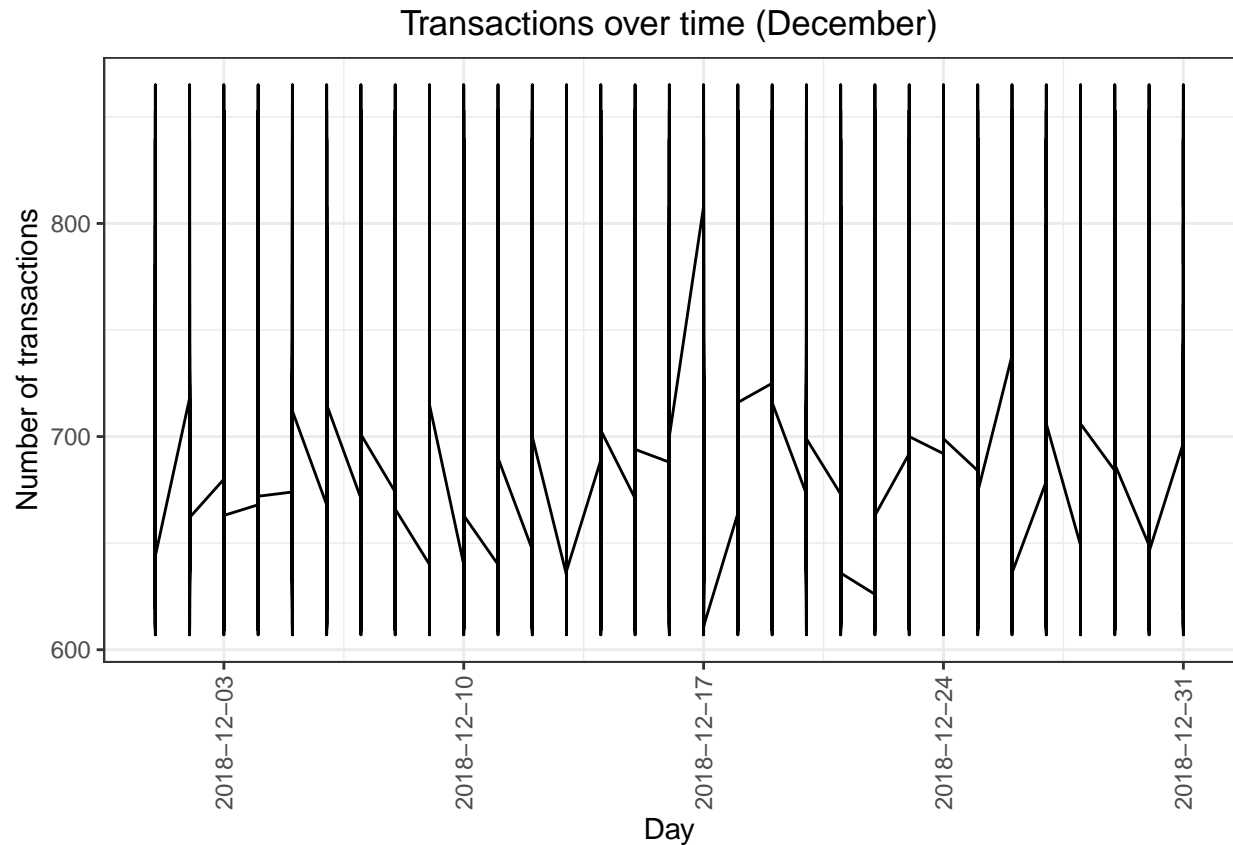


We can see that there is an increase in purchases in December and a break in late December. Let's zoom in on this.

```
#### Filter to December and look at individual days
# Over to you - recreate the chart above zoomed in to the relevant dates.
# Set the plot themes
theme_set(theme_bw())
theme_update(plot.title = element_text(hjust = 0.5))

# Plot transactions over time, zoomed in to December
ggplot(transactions_by_day_filled, aes(x = DATE, y = count)) +
  geom_line() +
  labs(x = "Day", y = "Number of transactions", title = "Transactions over time (December)") +
  scale_x_date(breaks = "1 week", limits = as.Date(c("2018-12-01", "2018-12-31"))) +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5))

## Warning: Removed 121576 rows containing missing values (`geom_line()`).
```



We can see that the increase in sales occurs in the lead-up to Christmas and that there are zero sales on Christmas day itself. This is due to shops being closed on Christmas day. Now that we are satisfied that the data no longer has outliers, we can move on to creating other features such as brand of chips or pack size from `PROD_NAME`. We will start with pack size.

```
#### Pack size
#### We can work this out by taking the digits that are in PROD_NAME
transactionData[, PACK_SIZE := parse_number(PROD_NAME)]
#### Always check your output
#### Let's check if the pack sizes look sensible
transactionData[, .N, PACK_SIZE][order(PACK_SIZE)]
```

##	PACK_SIZE	N
## 1:	70	1507
## 2:	90	3008
## 3:	110	22387
## 4:	125	1454
## 5:	134	25102
## 6:	135	3257
## 7:	150	40203
## 8:	160	2970
## 9:	165	15297
## 10:	170	19983
## 11:	175	66390
## 12:	180	1468
## 13:	190	2995
## 14:	200	4473
## 15:	210	6272




```
## 16:      220  1564
## 17:      250  3169
## 18:      270  6285
## 19:      330 12540
## 20:      380  6416
```

The largest size is 380g and the smallest size is 70g - seems sensible!

```
#### Let's plot a histogram of PACK_SIZE since we know that it is a categorical variable and not a cont.
# Over to you! Plot a histogram showing the number of transactions by pack size.
library(ggplot2)
```

```
# Plot histogram of transactions by pack size
ggplot(transactionData, aes(x = PACK_SIZE)) +
  geom_histogram(binwidth = 1) +
  labs(x = "Pack Size", y = "Number of Transactions", title = "Number of Transactions by Pack Size")
```



Pack sizes created look reasonable. Now to create brands, we can use the first word in PROD\_NAME to work out the brand name...

```
#### Brands
# Over to you! Create a column which contains the brand of the product, by extracting it from the product name.
#### Checking brands
# Over to you! Check the results look reasonable.
### Brands
library(stringr)

transactionData$brand_name <- str_to_title(word(transactionData$PROD_NAME, 1))
```

```
# view the first 10 rows
head(transactionData$brand_name, 10)

## [1] "Natural" "Ccs"      "Smiths"  "Smiths"  "Kettle"  "Smiths"  "Grain"
## [8] "Doritos" "Grain"    "Smiths"

unique(transactionData$brand_name)

## [1] "Natural"      "Ccs"          "Smiths"       "Kettle"       "Grain"
## [6] "Doritos"      "Twisties"     "Ww"           "Thins"        "Burger"
## [11] "Ncc"          "Cheezels"     "Infzns"       "Red"          "Pringles"
## [16] "Dorito"       "Infuzions"    "Smith"        "Grnwves"      "Tyrrells"
## [21] "Cobs"         "French"       "Rrd"          "Tostitos"     "Cheetos"
## [26] "Woolworths"  "Snbts"       "Sunbites"
```

Some of the brand names look like they are of the same brands - such as RED and RRD, which are both Red Rock Deli chips. Let's combine these together.

```
#### Clean brand names
transactionData[brand_name == "Red", brand_name := "Rrd"]
transactionData[brand_name == "Dorito", brand_name := "Doritos"]
# Over to you! Add any additional brand adjustments you think may be required.
#### Check again
# Over to you! Check the results look reasonable.
```

## Examining customer data

Now that we are happy with the transaction dataset, let's have a look at the customer dataset.

```
#### Examining customer data
# Over to you! Do some basic summaries of the dataset, including distributions of any key columns.
summary(customerData)
```

```
##  LYLTY_CARD_NBR      LIFESTAGE      PREMIUM_CUSTOMER
##  Min.   : 1000      Length:72637      Length:72637
##  1st Qu.: 66202      Class :character  Class :character
##  Median : 134040      Mode  :character  Mode  :character
##  Mean   : 136186
##  3rd Qu.: 203375
##  Max.   :2373711
```

```
tibble(customerData)
```

```
## # A tibble: 72,637 x 3
##   LYLTY_CARD_NBR LIFESTAGE      PREMIUM_CUSTOMER
##   <int> <chr>      <chr>
## 1      1000 YOUNG SINGLES/COUPLES Premium
## 2      1002 YOUNG SINGLES/COUPLES Mainstream
## 3      1003 YOUNG FAMILIES      Budget
## 4      1004 OLDER SINGLES/COUPLES Mainstream
## 5      1005 MIDAGE SINGLES/COUPLES Mainstream
## 6      1007 YOUNG SINGLES/COUPLES Budget
## 7      1009 NEW FAMILIES      Premium
## 8      1010 YOUNG SINGLES/COUPLES Mainstream
## 9      1011 OLDER SINGLES/COUPLES Mainstream
## 10     1012 OLDER FAMILIES      Mainstream
## # i 72,627 more rows
```

```

# Group the data by words and count the frequency of occurrence
word_freq <- customerData %>%
  group_by(LIFESTAGE) %>%
  summarize(count = n())
# Sort the data in descending order of frequency
word_freq <- word_freq %>%
  arrange(desc(count))
# View the top 10 most frequent words
head(word_freq, 10)

```

```

## # A tibble: 7 x 2
##   LIFESTAGE      count
##   <chr>         <int>
## 1 RETIREES      14805
## 2 OLDER SINGLES/COUPLES 14609
## 3 YOUNG SINGLES/COUPLES 14441
## 4 OLDER FAMILIES      9780
## 5 YOUNG FAMILIES      9178
## 6 MIDAGE SINGLES/COUPLES 7275
## 7 NEW FAMILIES       2549

```

```

# Group the data by words and count the frequency of occurrence
word_freq <- customerData %>%
  group_by(PREMIUM_CUSTOMER) %>%
  summarize(count = n())
# Sort the data in descending order of frequency
word_freq <- word_freq %>%
  arrange(desc(count))
# View the top 10 most frequent words
head(word_freq, 10)

```

```

## # A tibble: 3 x 2
##   PREMIUM_CUSTOMER count
##   <chr>             <int>
## 1 Mainstream       29245
## 2 Budget           24470
## 3 Premium          18922

```

```

#### Merge transaction data to customer data
data <- merge(transactionData, customerData, all.x = TRUE)

```

As the number of rows in `data` is the same as that of `transactionData`, we can be sure that no duplicates were created. This is because we created `data` by setting `all.x = TRUE` (in other words, a left join) which means take all the rows in `transactionData` and find rows with matching values in shared columns and then joining the details in these rows to the `x` or the first mentioned table. Page 7 20200128\_InsideSh-erpa\_Task1\_DraftSolutions - Template (1).Rmd Let's also check if some customers were not matched on by checking for nulls.

```

# Over to you! See if any transactions did not have a matched customer.
# Find transactions without a matched customer
unmatched_transactions <- anti_join(transactionData, customerData, by = "LYLTY_CARD_NBR")

# Check if there are any unmatched transactions
if (nrow(unmatched_transactions) == 0) {
  cat("All transactions have a matched customer.\n")
} else {

```

```
cat(paste0(nrow(unmatched_transactions), " transactions did not have a matched customer.\n"))
}
```

```
## All transactions have a matched customer.
```

Great, there are no nulls! So all our customers in the transaction data has been accounted for in the customer dataset. Note that if you are continuing with Task 2, you may want to retain this dataset which you can write out as a csv

```
fwrite(data, paste0(filePath, "QVI_data.csv"))
```

Data exploration is now complete! *## Data analysis on customer segments* Now that the data is ready for analysis, we can define some metrics of interest to the client: - Who spends the most on chips (total sales), describing customers by lifestage and how premium their general purchasing behaviour is - How many customers are in each segment - How many chips are bought per customer by segment - What's the average chip price by customer segment We could also ask our data team for more information. Examples are: - The customer's total spend over the period and total spend for each transaction to understand what proportion of their grocery spend is on chips - Proportion of customers in each customer segment overall to compare against the mix of customers who purchase chips Let's start with calculating total sales by LIFESTAGE and PREMIUM\_CUSTOMER and plotting the split by these segments to describe which customer segment contribute most to chip sales.

```
#### Total sales by LIFESTAGE and PREMIUM_CUSTOMER
```

```
# Over to you! Calculate the summary of sales by those dimensions and create a plot.
```

```
sales_summary <- data %>%
  group_by(LIFESTAGE, PREMIUM_CUSTOMER) %>%
  summarise(total_sales = sum(TOT_SALES)) %>%
  arrange(desc(total_sales))
```

```
## `summarise()` has grouped output by 'LIFESTAGE'. You can override using the
## `.groups` argument.
```

```
# View the summary
```

```
sales_summary
```

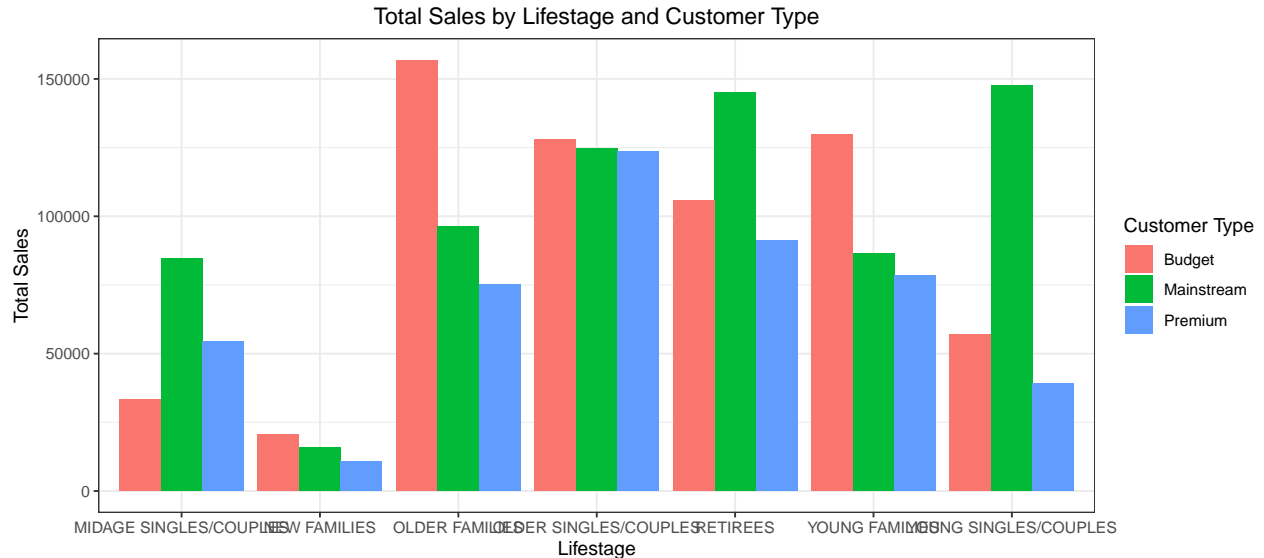
```
## # A tibble: 21 x 3
## # Groups:   LIFESTAGE [7]
##   LIFESTAGE      PREMIUM_CUSTOMER total_sales
##   <chr>          <chr>          <dbl>
## 1 OLDER FAMILIES Budget          156864.
## 2 YOUNG SINGLES/COUPLES Mainstream      147582.
## 3 RETIREES       Mainstream      145169.
## 4 YOUNG FAMILIES Budget          129718.
## 5 OLDER SINGLES/COUPLES Budget          127834.
## 6 OLDER SINGLES/COUPLES Mainstream      124648.
## 7 OLDER SINGLES/COUPLES Premium          123538.
## 8 RETIREES       Budget          105916.
## 9 OLDER FAMILIES Mainstream       96414.
## 10 RETIREES      Premium           91297.
## # i 11 more rows
```

```
library(ggplot2)
```

```
# Create the plot
```

```
ggplot(sales_summary, aes(x = LIFESTAGE, y = total_sales, fill = PREMIUM_CUSTOMER)) +
  geom_col(position = "dodge") +
  labs(title = "Total Sales by Lifestage and Customer Type",
```

```
x = "Lifestage",
y = "Total Sales",
fill = "Customer Type") +
theme(plot.title = element_text(hjust = 0.5))
```



Page 8 20200128\_InsideSherpa\_Task1\_DraftSolutions - Template (1).Rmd Sales are coming mainly from Budget - older families, Mainstream - young singles/couples, and Mainstream - retirees Let's see if the higher sales are due to there being more customers who buy chips.

```
#### Number of customers by LIFESTAGE and PREMIUM_CUSTOMER
# Over to you! Calculate the summary of number of customers by those dimensions and create a plot.
customer_count <- data %>%
  count(LIFESTAGE, PREMIUM_CUSTOMER, name = "customer_count") %>%
  arrange(desc(customer_count))

# View the results
customer_count
```


##	LIFESTAGE	PREMIUM_CUSTOMER	customer_count
## 1:	OLDER FAMILIES	Budget	21514
## 2:	RETIREES	Mainstream	19970
## 3:	YOUNG SINGLES/COUPLES	Mainstream	19544
## 4:	YOUNG FAMILIES	Budget	17763
## 5:	OLDER SINGLES/COUPLES	Budget	17172
## 6:	OLDER SINGLES/COUPLES	Mainstream	17061
## 7:	OLDER SINGLES/COUPLES	Premium	16560
## 8:	RETIREES	Budget	14225
## 9:	OLDER FAMILIES	Mainstream	13241
## 10:	RETIREES	Premium	12236
## 11:	YOUNG FAMILIES	Mainstream	11947
## 12:	MIDAGE SINGLES/COUPLES	Mainstream	11095
## 13:	YOUNG FAMILIES	Premium	10784
## 14:	OLDER FAMILIES	Premium	10403
## 15:	YOUNG SINGLES/COUPLES	Budget	8573
## 16:	MIDAGE SINGLES/COUPLES	Premium	7612
## 17:	YOUNG SINGLES/COUPLES	Premium	5852

```
## 18: MIDAGE SINGLES/COUPLES      Budget      4691
## 19:      NEW FAMILIES           Budget      2824
## 20:      NEW FAMILIES           Mainstream   2185
## 21:      NEW FAMILIES           Premium      1488
##                                LIFESTAGE PREMIUM_CUSTOMER customer_count
```

```
library(ggplot2)
```

```
# Create the plot
```

```
ggplot(customer_count, aes(x = LIFESTAGE, y = customer_count, fill = PREMIUM_CUSTOMER)) +
  geom_col(position = "dodge") +
  labs(title = "Customer Count by Lifestage and Customer Type",
       x = "Lifestage",
       y = "Customer Count",
       fill = "Customer Type") +
  theme(plot.title = element_text(hjust = 0.5))
```



There are more Mainstream - young singles/couples and Mainstream - retirees who buy chips. This contributes to there being more sales to these customer segments but this is not a major driver for the Budget - Older families segment. Higher sales may also be driven by more units of chips being bought per customer. Let's have a look at this next.

```
#### Average number of units per customer by LIFESTAGE and PREMIUM_CUSTOMER
```

```
# Over to you! Calculate and plot the average number of units per customer by those two dimensions.
```

```
prod_qty_avg <- data %>%
  group_by(LIFESTAGE, PREMIUM_CUSTOMER) %>%
  summarize(avg_qty = mean(PROD_QTY))
```

```
## `summarise()` has grouped output by 'LIFESTAGE'. You can override using the
## `.groups` argument.
```

```
# View the summary
```

```
prod_qty_avg
```


```
## # A tibble: 21 x 3
## # Groups:   LIFESTAGE [7]
##   LIFESTAGE      PREMIUM_CUSTOMER avg_qty
```

```
##      <chr>                <chr>                <dbl>
## 1 MIDAGE SINGLES/COUPLES Budget                1.89
## 2 MIDAGE SINGLES/COUPLES Mainstream            1.91
## 3 MIDAGE SINGLES/COUPLES Premium              1.89
## 4 NEW FAMILIES           Budget                1.86
## 5 NEW FAMILIES           Mainstream            1.86
## 6 NEW FAMILIES           Premium              1.86
## 7 OLDER FAMILIES         Budget                1.95
## 8 OLDER FAMILIES         Mainstream            1.95
## 9 OLDER FAMILIES         Premium              1.95
## 10 OLDER SINGLES/COUPLES Budget                1.91
## # i 11 more rows
```

```
library(ggplot2)
```

```
# Create the plot
```

```
ggplot(prod_qty_avg, aes(x = LIFESTAGE, y = avg_qty, fill = PREMIUM_CUSTOMER)) +
  geom_col(position = "dodge") +
  labs(x = "LIFESTAGE", y = "Average Quantity per Customer",
       title = "Average Number of Units per Customer by LIFESTAGE and PREMIUM_CUSTOMER") +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5))
```



Older families and young families in general buy more chips per customer Let's also investigate the average price per unit chips bought for each customer segment as this is also a driver of total sales.

```
#### Average price per unit by LIFESTAGE and PREMIUM_CUSTOMER
```

```
# Over to you! Calculate and plot the average price per unit sold (average sale price) by those two cus
```

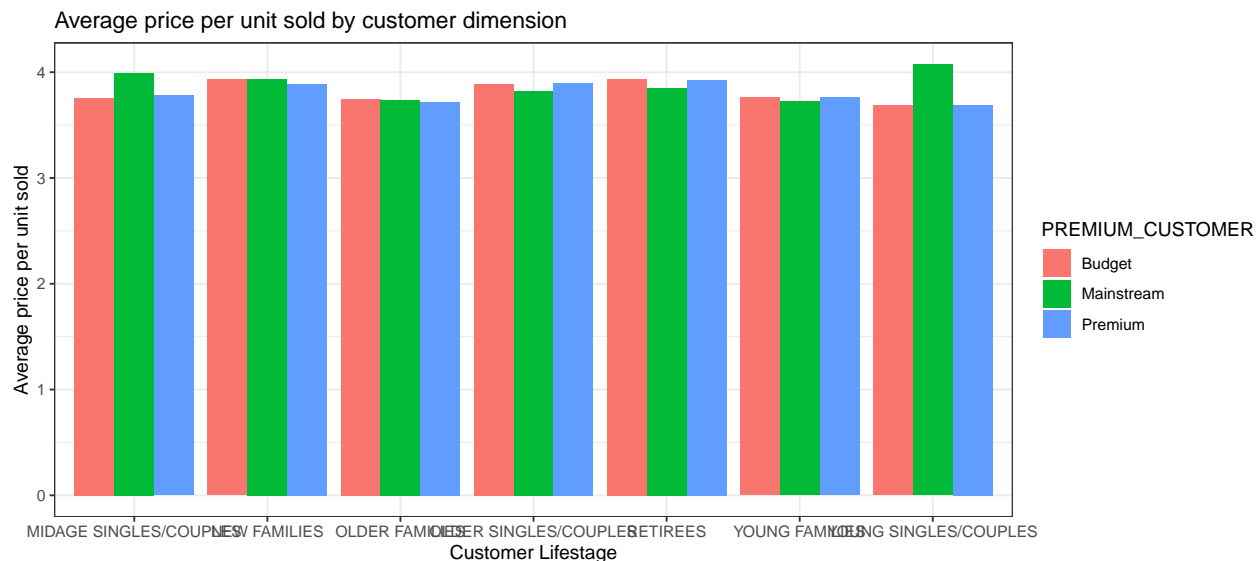
```
price_per_unit <- data %>%
  group_by(LIFESTAGE, PREMIUM_CUSTOMER) %>%
  summarize(avg_price_per_unit = sum(TOT_SALES) / sum(PROD_QTY))
```

```
## `summarise()` has grouped output by 'LIFESTAGE'. You can override using the
## `.groups` argument.
```

```
# Plot average price per unit sold
```

```
ggplot(price_per_unit, aes(x = LIFESTAGE, y = avg_price_per_unit, fill = PREMIUM_CUSTOMER)) +
  geom_bar(stat = "identity", position = "dodge") +
```

```
labs(x = "Customer Lifestage", y = "Average price per unit sold", title = "Average price per unit sold", theme_bw())
```



Mainstream midage and young singles and couples are more willing to pay more per packet of chips compared to their budget and premium counterparts. This may be due to premium shoppers being more likely to buy healthy snacks and when they buy chips, this is mainly for entertainment purposes rather than their own consumption. This is also supported by there being fewer premium midage and young singles and couples buying chips compared to their mainstream counterparts. As the difference in average price per unit isn't large, we can check if this difference is statistically different.

```
#### Perform an independent t-test between mainstream vs premium and budget midage and
#### young singles and couples
# Over to you! Perform a t-test to see if the difference is significant.
#### Perform an independent t-test between mainstream vs premium and budget midage and
#### young singles and couples
```

```
# Filter the relevant data
```

```
t_test_data <- price_per_unit %>%
  filter((LIFESTAGE %in% c("MIDAGE SINGLES/COUPLES", "YOUNG SINGLES/COUPLES")) &
    (PREMIUM_CUSTOMER %in% c("Budget", "Premium"))) |
    (LIFESTAGE %in% c("MIDAGE SINGLES/COUPLES", "YOUNG SINGLES/COUPLES")) &
    (PREMIUM_CUSTOMER == "Mainstream"))
```

```
group1 <- t_test_data %>%
  filter(LIFESTAGE %in% c("MIDAGE SINGLES/COUPLES", "YOUNG SINGLES/COUPLES") &
    PREMIUM_CUSTOMER == "Mainstream") %>%
  pull(avg_price_per_unit)
```

```
group2 <- t_test_data %>%
  filter(LIFESTAGE %in% c("MIDAGE SINGLES/COUPLES", "YOUNG SINGLES/COUPLES") &
    PREMIUM_CUSTOMER %in% c("Budget", "Premium")) %>%
  pull(avg_price_per_unit)
```



```
# Perform independent t-test
t_test_result <- t.test(group1, group2)
```

```
# View the test results
t_test_result
```

```
##
## Welch Two Sample t-test
##
## data: group1 and group2
## t = 6.6358, df = 1.7353, p-value = 0.03113
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## 0.07556858 0.53647998
## sample estimates:
## mean of x mean of y
## 4.034246 3.728222
```

The t-test results in a p-value of XXXXXXXX, i.e. the unit price for mainstream, young and mid-age singles and couples [ARE / ARE NOT] significantly higher than that of budget or premium, young and midage singles and couples. *## Deep dive into specific customer segments for insights* We have found quite a few interesting insights that we can dive deeper into. We might want to target customer segments that contribute the most to sales to retain them or further increase sales. Let's look at Mainstream - young singles/couples. For instance, let's find out if they tend to buy a particular brand of chips.

```
#### Deep dive into Mainstream, young singles/couples
```

```
# Over to you! Work out of there are brands that these two customer segments prefer more than others. Y
```

```
#install.packages("arules")
```

```
#install.packages("tidyverse")
```

```
#library(tidyverse)
```

```
#library(arules)
```

```
#library(dplyr)
```

```
# Create a subset of the data with only Mainstream and Young Singles/Couples customers
```

```
#customer_subset <- data %>%
```

```
## filter(LIFESTAGE %in% c("YOUNG SINGLES/COUPLES", "MIDAGE SINGLES/COUPLES") & PREMIUM_CUSTOMER == "M
```

```
# Create a transaction matrix
```

```
#transactions <- customer_subset %>%
```

```
# select(LYLTY_CARD_NBR, brand_name) %>%
```

```
# distinct() %>%
```

```
# group_by(LYLTY_CARD_NBR) %>%
```

```
# summarize(items = list(brand_name))
```

```
# Convert the transaction matrix to a binary matrix
```

```
#binary_matrix <- as.data.frame.matrix(table(transactions$LYLTY_CARD_NBR, #unlist(transactions$items)))
```

```
# Convert binary_matrix to transactions
```

```
#transactions <- as(transpose(binary_matrix), "transactions")
```

```
# Run the apriori algorithm with a minimum support of 0.01 and confidence of 0.5
```

```
#rules <- apriori(transactions, parameter = list(supp = 0.01, conf = 0.5))
```

```
# Inspect the rules  
#inspect(rules)
```

```
#### Preferred pack size compared to the rest of the population  
# Over to you! Do the same for pack size.
```